

MAPREDUCE 실습

학습 목표

- Python 으로 mapper와 reducer를 구현해 확인해 본다.
- cloudera 이용하여 mapreduce를 수행해 본다.

MapReduce 잡 실행

환경 :

Centos 6.3
Hadoop 2.x
Cloudera
vmware

01 Hadoop 버전 확인

\$ hadoop version

[training@localhost ~]\$ **hadoop version**

Hadoop 2.0.0-cdh4.1.1

Subversion <file:///data/1/jenkins/workspace/generic-package-centos32-6/topdir/BUILD/hadoop-2.0.0-cdh4.1.1/src/hadoop-common-project/hadoop-common> -r 581959ba23e4af85afd8db98b7687662fe9c5f20

Compiled by jenkins on Tue Oct 16 11:07:59 PDT 2012

From source with checksum 95f5c7f30b4030f1f327758e7b2bd61f

02 myinput 디렉터리 생성

hadoop fs -mkdir myinput

hadoop fs -ls /user/training [user 명이 training일 경우]

02 myinput에 purchases.txt를 넣기 (430만줄, 201MB)

디렉터리 이동 후, 파일을 myinput 으로 넣기

[cd /home/training/udacity_training/data](#)

[hadoop fs -put purchases.txt myinput](#)

[training@localhost data]\$ **hadoop fs -ls myinput**

Found 1 items

```
-rw-r--r-- 1 training supergroup 211312924 2019-02-19 09:52 myinput/purchases.txt
```

1	2	3	4	5	6
2012-01-01	09:00	San Jose	Men's Clothing	214,05	Amex↓
2012-01-01	09:00	Fort Worth	Women's Clothing	153,57	Visa↓
2012-01-01	09:00	San Diego	Music	66,08	Cash↓
2012-01-01	09:00	Pittsburgh	Pet Supplies	493,51	Discover↓
2012-01-01	09:00	Omaha	Children's Clothing	235,63	MasterCard↓
2012-01-01	09:00	Stockton	Men's Clothing	247,18	MasterCard↓
2012-01-01	09:00	Austin	Cameras	379,6	Visa↓
2012-01-01	09:00	New York	Consumer Electronics	296,8	Cash↓
2012-01-01	09:00	Corpus Christi	Toys	25,38	Discover↓
2012-01-01	09:00	Fort Worth	Toys	213,88	Visa↓
2012-01-01	09:00	Las Vegas	Video Games	53,26	Visa↓
2012-01-01	09:00	Newark	Video Games	39,75	Cash↓
2012-01-01	09:00	Austin	Cameras	469,63	MasterCard↓

03 mapper and reducer 실행

(아래 명령줄은 한줄이다. 길기에 두줄로 분리시켜둠. 실제 명령넣을 때는 한줄로 넣어야 함)

[사용법]

```
hadoop jar, <a path to a jar> -mapper <mapper> -reducer <reducer> -file <맵퍼 코드 파일> -file <리듀서 코드 파일>
-input <HDFS의 입력디렉터리> -output <출력 데이터를 쓸 리듀서의 출력 디렉터리>
```

해당 소스 디렉터리로 이동 후, 명령 실행

디렉터리 이동

```
cd /home/training/udacity_training/code
```

mapper와 reducer 실행

```
hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.0.0-mr1-cdh4.1.1.jar -mapper mapper.py -reducer reducer.py -
file mapper.py -file reducer.py -input myinput -output joboutput
```

(* mapper.py, reducer.py 코드가 있는 곳에서 실행)

04 결과 확인

결과 확인

```
hadoop fs -ls
```

joboutput 폴더의 리스트 확인

hadoop fs -ls joboutput

확인 파일

... joboutput/_SUCCESS => 성공적인 수행
.... joboutput/_logs => job이 실행되는 동안 일어났던 정보 로그를 포함.
.... joboutput/part-00000 => 우리가 수행했던 Job(잡) 하나의 리듀서로부터의 결과 파일

[실행 결과]

```
[training@localhost code]$ hadoop fs -ls
```

```
Found 2 items
```

```
drwxr-xr-x - training supergroup      0 2018-05-07 06:31 joboutput
```

```
drwxr-xr-x - training supergroup      0 2018-05-07 06:22 myinput
```

```
[training@localhost code]$ hadoop fs -ls joboutput
```

```
Found 3 items
```

```
-rw-r--r--  1 training supergroup      0 2018-05-07 06:31 joboutput/_SUCCESS
```

```
drwxr-xr-x - training supergroup      0 2018-05-07 06:27 joboutput/_logs
```

```
-rw-r--r--  1 training supergroup 2296 2018-05-07 06:31 joboutput/part-00000
```

생성된 파일의 내용 보기 (less : Enter 키 입력 한줄씩 출력, Space bar 한 화면씩 출력)

```
hadoop fs -cat joboutput/part-00000 | less
```

HDFS의 생성된 결과를 내 pc의 txt파일로 가져오기

```
hadoop fs -get joboutput/part-00000 mylocfile.txt
```

```
[training@localhost code]$ head -n 20 mylocfile.txt
```

```
Albuquerque 10052311.42
```

```
Anaheim 10076416.36
```

```
Anchorage 9933500.4
```

```
Arlington 10072207.97
```

```
Atlanta 9997146.7
```

```
Aurora 9992970.92
```

```
Austin 10057158.9
```

```
Bakersfield 10031208.92
```

```
Baltimore 10096521.45
```

```
Baton Rouge 10131273.23
```

```
Birmingham 10076606.52
```

```
Boise 10039166.74
```

```
Boston 10039473.28
```

```
Buffalo 10001941.19
```

```
Chandler 9919559.86
```

```
Charlotte 10112531.34
```

Chesapeake 10038504.92
 Chicago 10062522.07
 Chula Vista 9974951.34
 Cincinnati 10139505.74

05 Hadoop job 명령 실행

shell에 의한 mapper와 reducer의 실행

`hs mapper.py reducer.py myinput joboutput`

[실행 결과] joboutput 가 있어, 디렉터리가 존재한다고 함.

`hs mapper.py reducer.py myinput newoutputdir`

[결과]

[training@localhost code]\$ **hs mapper.py reducer.py myinput joboutput**

packageJobJar: [mapper.py, reducer.py, /tmp/hadoop-training/hadoop-unjar9181741705364652401/] [] /tmp/streamjob8795202652725235274.jar tmpDir=null

19/02/19 10:06:54 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should implement Tool for the same.

19/02/19 10:06:55 INFO mapred.JobClient: Cleaning up the staging area <https://0.0.0.0:8020/var/lib/hadoop->

https://0.0.0.0:8020/cache/mapred/mapred/staging/training/.staging/job_201902190936_0002

19/02/19 10:06:55 ERROR security.UserGroupInformation: PrivilegedActionException as:training (auth:SIMPLE) cause:org.apache.hadoop.mapred.FileAlreadyExistsException:

Output directory <https://0.0.0.0:8020/user/training/joboutput> already exists

19/02/19 10:06:55 ERROR streaming.StreamJob: Error launching job , Output path already exists : Output directory <https://0.0.0.0:8020/user/training/joboutput> already exists

Streaming Command Failed!

참고 자료 hs

hs 스크립트 소스 코드 확인(

```
run_mapreduce ()
{
  hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.0.0-mr1-cdh4.1.1.jar -mapper $1 -reducer $2 -file $1 -file $2 -input $3 -output $4
}
```

[실행결과]

[training@localhost data]\$ `hadoop fs -mkdir myinput`

[training@localhost data]\$ `hadoop fs -put purchases.txt myinput`

[training@localhost data]\$ `hadoop fs -ls myinput`

Found 1 items

-rw-r--r-- 1 training supergroup 211312924 2018-05-07 06:22 myinput/purchases.txt

```
[training@localhost code]$ hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.0.0-mr1-cdh4.1.1.jar -mapper mapper.py -reducer reducer.py -file mapper.py -file reducer.py -input myinput -output joboutput
```

```
packageJobJar: [mapper.py, reducer.py, /tmp/hadoop-training/hadoop-unjar4354191662551748265/] [] /tmp/streamjob389135337284658968.jar tmpDir=null
```

```
18/05/07 06:27:23 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should implement Tool for the same.
```

```
18/05/07 06:27:24 WARN snappy.LoadSnappy: Snappy native library is available
```

```
18/05/07 06:27:24 INFO snappy.LoadSnappy: Snappy native library loaded
```

```
18/05/07 06:27:24 INFO mapred.FileInputFormat: Total input paths to process : 1
```

```
18/05/07 06:27:25 INFO streaming.StreamJob: getLocalDirs(): [/var/lib/hadoop-hdfs/cache/training/mapred/local]
```

```
18/05/07 06:27:25 INFO streaming.StreamJob: Running job: job_201805070616_0001
```

```
18/05/07 06:27:25 INFO streaming.StreamJob: To kill this job, run:
```

```
18/05/07 06:27:25 INFO streaming.StreamJob: UNDEF/bin/hadoop job -Dmapred.job.tracker=0.0.0.0:8021 -kill job_201805070616_0001
```

```
18/05/07 06:27:25 INFO streaming.StreamJob: Tracking URL: http://0.0.0.0:50030/jobdetails.jsp?jobid=job\_201805070616\_0001
```

```
18/05/07 06:27:26 INFO streaming.StreamJob: map 0% reduce 0%
```

```
18/05/07 06:27:49 INFO streaming.StreamJob: map 3% reduce 0%
```

```
18/05/07 06:27:52 INFO streaming.StreamJob: map 5% reduce 0%
```

```
18/05/07 06:27:56 INFO streaming.StreamJob: map 9% reduce 0%
```

```
18/05/07 06:27:59 INFO streaming.StreamJob: map 11% reduce 0%
```

```
18/05/07 06:28:02 INFO streaming.StreamJob: map 12% reduce 0%
```

```
18/05/07 06:28:03 INFO streaming.StreamJob: map 13% reduce 0%
```

```
18/05/07 06:28:06 INFO streaming.StreamJob: map 15% reduce 0%
```

```
18/05/07 06:28:09 INFO streaming.StreamJob: map 18% reduce 0%
```

```
...
```

```
18/05/07 06:28:51 INFO streaming.StreamJob: map 49% reduce 0%
```

```
18/05/07 06:28:54 INFO streaming.StreamJob: map 50% reduce 0%
```

```
18/05/07 06:29:22 INFO streaming.StreamJob: map 53% reduce 8%
```

```
18/05/07 06:29:25 INFO streaming.StreamJob: map 58% reduce 8%
```

```
18/05/07 06:29:26 INFO streaming.StreamJob: map 59% reduce 17%
```

```
18/05/07 06:29:29 INFO streaming.StreamJob: map 68% reduce 17%
```

```
18/05/07 06:29:32 INFO streaming.StreamJob: map 78% reduce 17%
```

```
18/05/07 06:29:35 INFO streaming.StreamJob: map 79% reduce 17%
```

```
18/05/07 06:29:36 INFO streaming.StreamJob: map 80% reduce 17%
```

```
18/05/07 06:29:39 INFO streaming.StreamJob: map 81% reduce 25%
```

```
18/05/07 06:29:42 INFO streaming.StreamJob: map 83% reduce 25%
```

```
18/05/07 06:29:45 INFO streaming.StreamJob: map 86% reduce 25%
```

```
...
```

```
18/05/07 06:30:01 INFO streaming.StreamJob: map 96% reduce 25%
```

```
18/05/07 06:30:04 INFO streaming.StreamJob: map 99% reduce 25%
```

```
18/05/07 06:30:07 INFO streaming.StreamJob: map 100% reduce 25%
```

```
18/05/07 06:30:13 INFO streaming.StreamJob: map 100% reduce 33%
```

```
18/05/07 06:30:22 INFO streaming.StreamJob: map 100% reduce 67%
```

```
18/05/07 06:30:25 INFO streaming.StreamJob: map 100% reduce 69%
```

```
18/05/07 06:30:28 INFO streaming.StreamJob: map 100% reduce 71%
...
18/05/07 06:31:08 INFO streaming.StreamJob: map 100% reduce 93%
18/05/07 06:31:11 INFO streaming.StreamJob: map 100% reduce 95%
18/05/07 06:31:15 INFO streaming.StreamJob: map 100% reduce 97%
18/05/07 06:31:18 INFO streaming.StreamJob: map 100% reduce 98%
18/05/07 06:31:21 INFO streaming.StreamJob: map 100% reduce 100%
18/05/07 06:31:27 INFO streaming.StreamJob: Job complete: job_201805070616_0001
18/05/07 06:31:27 INFO streaming.StreamJob: Output: joboutput
```

```
[training@localhost code]$ hadoop fs -ls
```

```
Found 2 items
```

```
drwxr-xr-x - training supergroup      0 2018-05-07 06:31 joboutput
drwxr-xr-x - training supergroup      0 2018-05-07 06:22 myinput
```

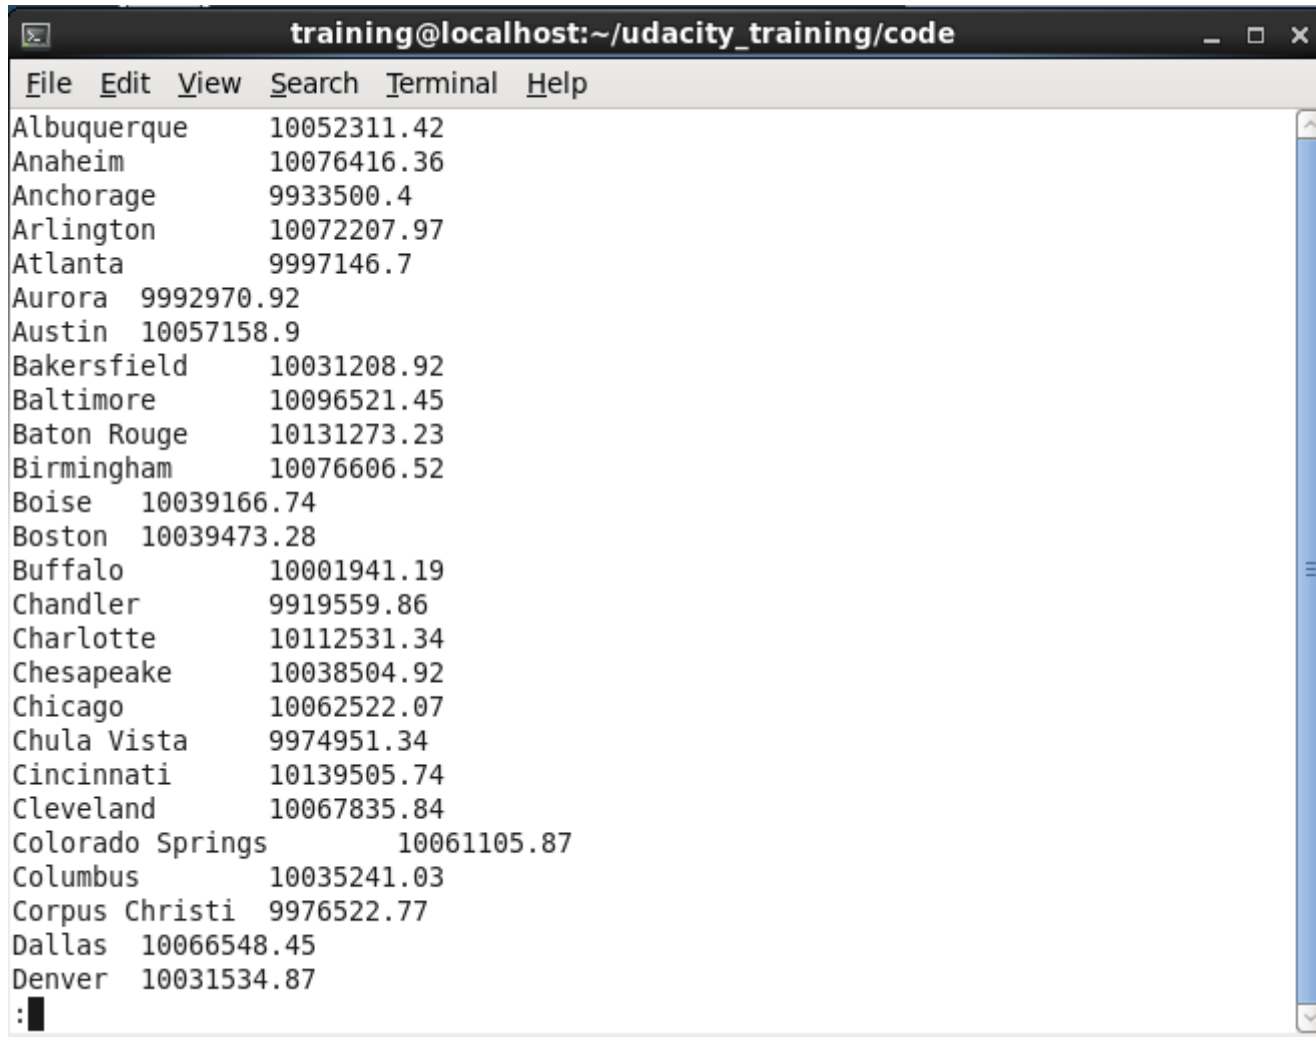
```
[training@localhost code]$ hadoop fs -ls joboutput
```

```
Found 3 items
```

```
-rw-r--r--  1 training supergroup      0 2018-05-07 06:31 joboutput/_SUCCESS
drwxr-xr-x - training supergroup      0 2018-05-07 06:27 joboutput/_logs
-rw-r--r--  1 training supergroup 2296 2018-05-07 06:31 joboutput/part-00000
```

```
hadoop fs -cat joboutput/part-00000 | less
```

[결과 파일]



```
training@localhost:~/udacity_training/code
File Edit View Search Terminal Help
Albuquerque 10052311.42
Anaheim 10076416.36
Anchorage 9933500.4
Arlington 10072207.97
Atlanta 9997146.7
Aurora 9992970.92
Austin 10057158.9
Bakersfield 10031208.92
Baltimore 10096521.45
Baton Rouge 10131273.23
Birmingham 10076606.52
Boise 10039166.74
Boston 10039473.28
Buffalo 10001941.19
Chandler 9919559.86
Charlotte 10112531.34
Chesapeake 10038504.92
Chicago 10062522.07
Chula Vista 9974951.34
Cincinnati 10139505.74
Cleveland 10067835.84
Colorado Springs 10061105.87
Columbus 10035241.03
Corpus Christi 9976522.77
Dallas 10066548.45
Denver 10031534.87
:
```

```
[training@localhost code]$ hadoop fs -get joboutput/part-00000 mylocfile.txt
```

```
[training@localhost code]$ ls -ltr
```

```
total 12
```

```
-rwxrwxr-x 1 training training 743 Sep  9 2013 reducer.py
```

```
-rwxrwxr-x 1 training training 424 Sep 10 2013 mapper.py
```

```
-rwxr-xr-x 1 training training 2296 May  7 06:38 mylocfile.txt
```

참고 python 소스 코드

```
# mapper.py
```

```
#!/usr/bin/python
```

```
# Format of each line is:
```

```
# date\ttime\tstore name\titem description\tcost\tmethod of payment
```

```
#
```

```
# We want elements 2 (store name) and 4 (cost)
```

```
# We need to write them out to standard output, separated by a tab
```

```
import sys
```

```
for line in sys.stdin:
```

```
    data = line.strip().split("\t")
```

```
    if len(data) == 6:
```

```
        date, time, store, item, cost, payment = data
```

```
        print "{0}\t{1}".format(store, cost)
```

```
# reducer.py
```

```
#!/usr/bin/python
```

```
import sys
```

```
salesTotal = 0
```

```
oldKey = None
```

```
# Loop around the data
```

```
# It will be in the format key\tval
```

```
# Where key is the store name, val is the sale amount
```

```
#
```

```
# All the sales for a particular store will be presented,
```

```
# then the key will change and we'll be dealing with the next store
```

```
for line in sys.stdin:
```

```
    data_mapped = line.strip().split("\t")
```

```
    if len(data_mapped) != 2:
```

```
        # Something has gone wrong. Skip this line.
```

```
        continue
```

```
    thisKey, thisSale = data_mapped
```



```
if oldKey and oldKey != thisKey:
    print oldKey, "\t", salesTotal
    oldKey = thisKey;
    salesTotal = 0

oldKey = thisKey
salesTotal += float(thisSale)

if oldKey != None:
    print oldKey, "\t", salesTotal
```

실습과제 1

위의 Hadoop의 MapReduce는 스토어별 총 합계를 구하였다.

(1) 이와 비슷하게 Mapper와 Reducer를 이용하여 item 별 총합계를 구해서 계산해 보자.
확인한 결과 화면을 캡처해 보자.

도전과제 1

(2) Item별 총 데이터 횟수를 구해보자.
(3) store별 총 데이터 횟수를 구해보자.

확인한 결과 화면을 캡처해 보자.